

Bayes' Rule for classification

Recall Bayes' rule:

$$P(A | B) = \frac{P(A)P(B | A)}{P(B)}$$

You might remember that each of these terms has a name:

- ▶ $P(A)$: the prior probability
- ▶ $P(A | B)$: the posterior probability
- ▶ $P(B | A)$: the likelihood
- ▶ $P(B)$: the marginal (total/overall) probability

In classification, “A” is a class label and “B” is a set of features.

Bayes' Rule for classification

Bayes's rule:

$$P(y = k | x) = \frac{P(y = k) \cdot P(x | y = k)}{P(x)}$$

$P(y = k)$ is the prior probability for class k . We usually get this from the raw class frequencies in the training data. For example:

```
table(fgl_train$type) %>% prop.table %>% round(3)
```

```
##
```

```
## WinF WinNF Veh Con Tabl Head
```

```
## 0.297 0.366 0.087 0.058 0.047 0.145
```

Bayes' Rule for classification

Bayes's rule:

$$P(y = k | x) = \frac{P(y = k) \cdot P(x | y = k)}{P(x)}$$

$P(x)$ is the marginal probability of observing feature vector x .
Notice it doesn't depend on k ! It's the same number for all classes.

Thus we usually write the posterior probabilities up to this constant of proportionality, without bothering to compute it:

$$P(y = k | x) \propto P(y = k) \cdot P(x | y = k)$$

(Note: often we do the actual computations on a log scale instead.)

Bayes' Rule for classification

Bayes's rule:

$$P(y = k \mid x) = \frac{P(y = k) \cdot P(x \mid y = k)}{P(x)}$$

The hard part is estimating the likelihood $P(x \mid y = k)$. In words: how likely is it that we would have observed feature vector x if the true class label were k ?

This is like regression in reverse!

Naive Bayes

Recall that $x = (x_1, x_2, \dots, x_p)$ is a vector of p features. The simplest strategy for estimating $P(x \mid y = k)$ is called “Naive Bayes.”

It’s “naive” because we make the simplifying assumption that *every feature x_j is independent* of all other features, conditional on the class labels:

$$\begin{aligned} P(x \mid y = k) &= P(x_1, x_2, \dots, x_p \mid y = k) \\ &= \prod_{j=1}^p P(x_j \mid y = k) \quad (\text{independence}) \end{aligned}$$

This simplifies the requirements of the problem: *just calculate the marginal distribution of the features*, i.e. $P(x_j \mid y = k)$ for all features j and classes k .

Naive Bayes: a small example

In `congress109.csv` we have data on all speeches given on the floor of the U.S. Congress during the 109th Congressional Session (January 3, 2005 to January 3, 2007).

Every row is a set of *phrase counts* associated with a single representative's speeches across the whole session. X_{ij} = number of times that rep i utter phrase j during a speech.

The target variable $y \in \mathcal{R}$, D is the party affiliation of the representative.

Naive Bayes: a small example

We'll focus on just a few phrases and famous politicians:

```
# read in data
congress109 = read.csv("../data/congress109.csv", header=TRUE, row.names=1)
congress109members = read.csv("../data/congress109members.csv", header=TRUE, row.names=1)
```

Focus on a few key phrases and a few famous pols:

```
X_small = dplyr::select(congress109, minimum.wage, war.terror, tax.relief, hurricane.katrina)
X_small[c('John McCain', 'Mike Pence', 'John Kerry', 'Edward Kennedy'),]
```

##	minimum.wage	war.terror	tax.relief	hurricane.katrina
## John McCain	0	27	0	14
## Mike Pence	0	12	1	11
## John Kerry	12	16	13	23
## Edward Kennedy	260	8	1	53

Naive Bayes: a small example

Let's look at these counts summed across all members in each party:

```
y = congress109members$party
```

```
# Sum phrase counts by party
```

```
R_rows = which(y == 'R')
```

```
D_rows = which(y == 'D')
```

```
colSums(X_small[R_rows,])
```

```
##      minimum.wage      war.terror      tax.relief hurricane.katrina  
##           294           604           497           717
```

```
colSums(X_small[D_rows,])
```

```
##      minimum.wage      war.terror      tax.relief hurricane.katrina  
##           767           237           176           1295
```

So we get the sense that some phrases are “more Republican” and some “more Democrat.”

Naive Bayes: a small example

To make this precise, let's build our Naive Bayes model for a Congressional speech:

- ▶ Imagine that every phrase uttered in a speech is a random sample from a “bag of phrases,” where each phrase has its own probability. (*This is the Naive Bayes assumption of independence.*)
- ▶ Here the bag consists of just four phrases: “minimum wage”, “war on terror”, “tax relief,” and “hurricane katrina”.
- ▶ Each class (R or D) has its own probability vector associated with the phrases in the bag.

Naive Bayes: a small example

We can estimate these probability vectors for each class from the phrase counts in the training data.

For Republicans:

```
probat_R = colSums(X_small[R_rows,])  
probat_R = probat_R/sum(probat_R)  
probat_R %>% round(3)
```

##	minimum.wage	war.terror	tax.relief	hurricane.katrina
##	0.139	0.286	0.235	0.339

And for Democrats:

```
probat_D = colSums(X_small[D_rows,])  
probat_D = probat_D/sum(probat_D)  
probat_D %>% round(3)
```

##	minimum.wage	war.terror	tax.relief	hurricane.katrina
##	0.310	0.096	0.071	0.523

Naive Bayes: a small example

Let's now look at some particular member of Congress and try to build the likelihood, $P(x | y)$, for his or her phrase counts

```
X_small['Sheila Jackson-Lee',]
```

```
##               minimum.wage war.terror tax.relief hurricane.katrina
## Sheila Jackson-Lee           11          15           3           66
```

Are Sheila Jackson-Lee's phrase counts $x = (11, 15, 3, 66)$ more likely under the Republican or Democrat probability vector?

Naive Bayes: a small example

Recall $P(x \mid y = R)$:

##	minimum.wage	war.terror	tax.relief	hurricane.katrina
##	0.1392	0.2860	0.2353	0.3395

Under this probability vector:

$$\begin{aligned}P(x \mid y = R) &= P(x_1 = 11 \mid y = R) \\&\quad \times P(x_2 = 15 \mid y = R) \\&\quad \times P(x_3 = 3 \mid y = R) \\&\quad \times P(x_4 = 66 \mid y = R) \\&= (0.1392)^{11} \cdot (0.2860)^{15} \cdot (0.2353)^3 \cdot (0.3395)^{66} \\&= 3.765 \times 10^{-51}\end{aligned}$$

Naive Bayes: a small example

Now recall $P(x \mid y = D)$:

##	minimum.wage	war.terror	tax.relief	hurricane.katrina
##	0.1392	0.2860	0.2353	0.3395

Under this probability vector:

$$\begin{aligned}P(x \mid y = D) &= P(x_1 = 11 \mid y = D) \\&\quad \times P(x_2 = 15 \mid y = D) \\&\quad \times P(x_3 = 3 \mid y = D) \\&\quad \times P(x_4 = 66 \mid y = D) \\&= (0.3099)^{11} \cdot (0.0958)^{15} \cdot (0.0711)^3 \cdot (0.5232)^{66} \\&= 1.293 \times 10^{-43}\end{aligned}$$

Naive Bayes: a small example

These numbers are tiny, so it's much safer to work on a log scale:

$$\log P(x \mid y = k) = \sum_{j=1}^p x_j \log p_j^{(k)}$$

where $p_j^{(k)}$ is the j th entry in the probability vector for class k .

```
x_try = X_small['Sheila Jackson-Lee',]  
sum(x_try * log(probhat_R))
```

```
## [1] -116.1083
```

```
sum(x_try * log(probhat_D))
```

```
## [1] -98.75633
```

Naive Bayes: a small example

Let's use Bayes' rule (posterior \propto prior times likelihood) to put this together with our prior, estimated using the empirical class frequencies:

```
table(y) %>% prop.table %>% round(3)
```

```
## y
##      D      I      R
## 0.457 0.004 0.539
```

So:

$$P(R \mid x) \propto 0.539 \cdot (3.765 \times 10^{-51})$$

and

$$P(D \mid x) \propto 0.457 \cdot (1.293 \times 10^{-43})$$

Naive Bayes: a small example

To actually calculate a posterior, we must turn this into a set of probabilities by normalizing, i.e. dividing by the sum across all classes:

$$P(D \mid x) = \frac{0.457 \cdot (1.293 \times 10^{-43})}{0.457 \cdot (1.293 \times 10^{-43}) + 0.539 \cdot (3.765 \times 10^{-51})} \\ \approx 1$$

So:

1. Our model thinks Sheila Jackson-Lee is a Democrat.
2. The data completely overwhelm the prior! This is often the case in Naive Bayes models.

Naive Bayes: a bigger example

Let's turn to `congress109_bayes.R` to see a larger example of Naive Bayes classification, where we fit our model with all 1000 phrase counts.

Naive Bayes: summary

- ▶ Works by directly modeling $P(x | y)$, versus $P(y | x)$.
- ▶ This **regression in reverse** only works because we assume that each feature in x is independent, given the class labels.
- ▶ Simple and easy to compute, and therefore scalable to very large data sets and classification problems.
- ▶ Unlike a logit model, it works even more with features P than examples N .
- ▶ Often too simple: the “naive” assumption of independence really is a drastic simplification.
- ▶ The resulting probabilities are useful for classification purposes, but often not believable as probabilities.
- ▶ Most useful when the features x are categorical variables (like phrase counts!) Very common in text analysis.